Lab 6: Autocomplete

You are not logged in.

If you are a current student, please Log In for full access to the web site.
Note that this link will take you to an external site (`https://shimmer.csail.mit.edu`) to authenticate, and then you will be redirected back to this page.

Table of Contents

1) Preparation

This lab assumes you have Python 3.7 or later installed on your machine.

The following file contains code and other resources as a starting point for this lab: `lab6.zip`

Most of your changes should be made to `lab.py`, which you will submit at the end of this lab. Importantly, you should not add any imports to the file.

This lab is worth a total of 4 points. Your score for the lab is based on:

- correctly answering the questions on this page (0.5 points)
- passing the test cases from `test.py` under the time limit (1.5 points), and
- a brief "checkoff" conversation with a staff member to share your code, including discussion of the bullets below and a review of your code's clarity/style (2 points).

For this lab, you will only receive credit for a test case if it runs to completion under the time limit on the server.

The questions on this page (including your code submission) are due at 4pm on Friday, Oct 25. However, you are strongly encouraged to read sections 1-3, and perhaps to get started with trying to implement some of the methods, before lecture on Oct 21.

2) Introduction

Type "aren't you" into a search engine and you'll get a handful of search suggestions, ranging from "aren't you clever?" to "aren't you a little short for a stormtrooper?". If you've ever done a web search, you've probably seen an autocompletion — a handy list of words that pops up under your search, guessing at what you were about to type.

Search engines aren't the only place you'll find this mechanism. For example, cell phones use autocomplete/autocorrect to predict/correct words. Some IDEs (integrated development environments — used for coding and software development) use

autocomplete to make the process of coding more efficient by offering suggestions for completing long function or variable names.

In this lab, we are going to implement our own version of an autocomplete/autocorrect engine using a tree structure called a *trie*, as described in this document.

The lab will ask you first to create a class to represent a generic trie data structure. You will then use the trie to write your own autocomplete and autocorrect, as well as a mechanism for searching.

2.1) The Trie Data Structure

A trie[1], also known as a prefix tree, is a type of search tree that stores an associative array (a mapping from keys to values). In a trie, the keys are always ordered sequences. The trie stores keys organized by their prefixes (their first characters), with longer prefixes given by successive levels of the trie. Each node optionally contains a value to be associated with that node's prefix.

As an example, consider a trie constructed as follows:

```
t = Trie()
t['bat'] = True
t['bar'] = True
t['bark'] = True
```

This trie would look like the following (Fig. 1):

Fig. 1

One important thing to notice is that the keys associated with each node are not actually stored in the nodes themselves. Rather, they are stored in the edges connecting the nodes.

We'll start by implementing a class called `Trie` to represent tries in Python. This class will include facilities for adding, deleting, modifying, and retrieving all key/value pairs. For example, consider:

```
>>> t = Trie()
>>> t['bat'] = True
>>> t['bar'] = True
>>> t['bark'] = True
>>>
>>> t['bat']
True
>>> t['something']
Traceback (most recent call last):
    ...
KeyError
>>>
>>> t['bark'] = 20
>>> t['bark']
20
>>>
>>> for i in t:
>>>     print(i)

('bat', True)
('bar', True)
('bark', 20)
>>>
>>> del t['bar']
>>>
>>> for i in t:
>>>     print(i)

('bat', True)
('bark', 20)
```

However, we are not limited to using only strings. Your `Trie` structure should (eventually) also support using tuples as keys, for example:

```
>>> t = Trie()
>>> t[(2, )] = 'cat'
>>> t[(1, 0, 0)] 'dog'
>>> t[1, 0, 1)] = 'ferret'
>>> t[(1, 0, 1, 80)] = 'tomato'
>>>
>>> t[(1, 0)]
Traceback (most recent call last):
    ...
KeyError
>>> t[(1, 0, 0)]
'dog'
>>> for i in t:
>>>     print(i)

((2,), 'cat')
((1, 0, 0), 'dog')
((1, 0, 1), 'ferret')
((1, 0, 1, 80), 'tomato')
```

Note that, in terms of functionality, the `Trie` class will have a lot in common with a Python dictionary. However, the representation we're using "under the hood" has some nice features that make it well-suited for tasks (like autocompletion) that use prefix-based lookups.

3) Trie class and basic methods

In `lab.py`, you are responsible for implementing the `Trie` class, which should support the following methods.

Hint: you may wish to make sure everything is working for only `str` keys first, and then expand to make things work for keys that are `tuples`, rather than trying to implement both right from the start.

__init__(*self*)

Initialize self to be an object with exactly three instance variables:

- `value`, the value associated with the sequence ending at this node. Initial value is `None` (we will assume that a value of `None` means that a given key has no value associated with it, not that the value `None` is associated with it).
- `children`, a dictionary mapping single-element sequences (either length-1 strings, or length-1 tuples) to another trie node, i.e., the next level of the trie hierarchy (tries are a recursive data structure). Initial value is an empty dictionary.
- `type`, some way to keep track of the type of the keys (**without explicitly storing the entire keys themselves**). The exact choice of representation is up to you. This attribute should be set to `None` when the instance is first created, and it should be updated to reflect the type of the keys when the first element is added. You may assume that all keys in a given `Trie` instance are of the same type.

__getitem__(*self, key*)

Return the value associated with the given key. This is a special method name used by Python to implement subscripting (indexing). For example, `x[k]` is translated by Python into `x.__getitem__(k)`.

It is expected that the trie node descended from the `self` trie corresponding to `key` will be located and the value associated with that node will be returned, or a `KeyError` will be raised if the key cannot be found in the trie. If the type of the key is not consistent with the expected type of keys for this trie, raise a `TypeError`. More information on raising errors is [here](#).

Examples (using the example trie from above):

- `t['bar']` should return `True`.
- `t['apple']` should raise a `KeyError` since the given key does not exist in the trie.
- `t['ba']` should also raise a `KeyError` since, even though the key `'ba'` is represented in the trie, it has no value associated with it.
- `t[1]` should raise a `TypeError` since the keys for this trie are expected to be strings, not integers.

__setitem__(*self, key, value*)

This is a special method name used by Python to implement subscript assignment. For example, `x[k] = v` is translated by Python into `x.__setitem__(k, v)`.

Add the given `key` to the trie, associating it with the given `value`. For the trie node that marks the end of the key, set that node's `value` attribute to be given `value` argument. This method doesn't return a value. If the type of the `key` is not consistent with the key type expected for the trie, a `TypeError` exception should be raised.

Examples (using the trie structure from the picture above):

- `t = Trie()` would create the root node of the example trie above.
- `t['bat'] = True` adds three nodes (representing the `'b'`, `'ba'`, and `'bat'` prefixes), and associates the value `True` with the node corresponding to `'bat'`.
- `t['bark'] = True` adds two new nodes for prefixes `'bar'` and `'bark'` shown on the bottom right of the trie, setting the value of the last node to `True`.
- `t['bar'] = True` doesn't add any nodes and only sets the value of the first node added above when inserting "bark" to `True`.
- `t[1] = True` raises a `TypeError` and does not make any change to the trie.

__delitem__(*self, key*)

Disassociate the given key from its value. This is a special method name used by Python to implement index deletion. For example, `del x[k]` is translated by Python into `x.__delitem__(k)`.

Examples (using the example trie from above):

- `del t["bar"]` should disassociate `"bar"` from its value in the trie, so that subsequent lookups of `t["bar"]` produce a `KeyError`.

For the purposes of this lab, you only need to do the bare minimum so that the key is no longer associated with a value (don't worry about extra memory usage after deleting the key).

__contains__(*self, key*)

Return `True` if `key` occurs and has a value other than `None` in the trie. `__contains__` is the special method name used by Python to implement the `in` operator. For example,

```
k in x
```

is translated to

```
x.__contains__(k)
```

Hint: At first glance, the code for this method might look very similar to some of the other methods above. Make good use of helper functions to avoid repetitious code!

Examples (using the example trie from above):

- `"ba" in t` returns `False` since that interior node has no value associated with it.
- `"bar" in t` returns `True` (not because the value associated with `"bar"` is `True`, but because `"bar"` has a value associated with it at all.
- `"barking" in t` returns `False` since `"barking"` can't be found in the trie.

__iter__(*self*)

A generator that yields `(key, value)` tuples for each key stored in the trie. The pairs can be produced in any order. `__iter__` is the special method name used by Python when it needs to iterate over a data object, i.e., the method invoked by the `iter()` built-in function. For example, the following Python code:

```
for key, val in t:
    print(key)
```

Is equivalent to writing:

```
for key, val in t.__iter__():
    print(key)
```

And will print all of the `(key, val)` tuples in `t`.

Hint: You'll want to build a recursive generator function that uses `yield` and `yield from` to produce the required sequence of values one at a time. See the Python generator docs and/or more docs on generator syntax. You can also review the lecture 6 and recitation 8 materials.

Examples (using the example trie from above):

- `list(t)` returns `[('bat', True), ('bar', True), ('bark', True)]`. Note that the `list` function has an internal `for` loop that uses `iter(t)` to iterate over each element of the sequence `t`.

4) Autocomplete

Now, let's implement our auto-complete engine!

We'll start with implementing autocompletion for words, and then we'll move to implementing autocompletion for sentences. As a start for either of these, we'll need a way to build up a `Trie` instance from a text document.

make_word_trie(text)

> `text` is a string containing a body of text. Return a `Trie` instance mapping _words_ in the text to the frequency with which they occur in the given piece of text.
>
> Note that we have provided a method called `tokenize_sentences` which will try to intelligently split a piece of text into individual sentences. **You should use this function rather than implementing your own.** The function takes in a single string and returns a list of strings, one for each sentence, where punctuation has been stripped out and the sentence consists only of words. Words within those sentences are sequences of characters separated by spaces.

make_phrase_trie(text)

> `text` is a string containing a body of text. Return a `Trie` instance mapping sentences (represented as tuples of words) to the frequency with which they occur in the given piece of text.

As a running example, we'll use the following trie (Fig. 2), which could have been created by calling `make_word_trie("bat bat bark bar")`:


Fig. 2

Once we have those trie representations, we are ready to go ahead and implement autocompletion! We'll implement autocompletion as a function described below:

autocomplete(*trie, prefix, max_count*=None)

`trie` is an instance of `Trie`, `prefix` is a string/tuple, `max_count` is an integer or `None`. Return a list of the `max_count` most-frequently-occurring keys that start with `prefix`. In the case of a tie, you may output any of the most-frequently-occurring keys. If there are fewer than `max_count` valid keys available starting with `prefix`, return only as many as there are. The returned list may be in any order. If `max_count` is not specified, your list should contain _all_ keys that start with `prefix`.

Return `[]` if `prefix` is not in the trie. Raise a `TypeError` if the given prefix has the wrong type.

Examples (using the example trie from above):

- `autocomplete(t, "ba", 1)` returns `['bat']`.
- `autocomplete(t, "ba", 2)` might return either `['bat', 'bark']`, `['bark', 'bat']`, `['bat', 'bar']`, or `['bar', 'bat']` since "bark" and "bar" occur with equal frequency.
- `autocomplete(t, "be", 1)` returns `[]`.

Your implementation should be agnostic to the type of its inputs (i.e., it should work both on tries/prefixes that are either strings or tuples). Write a few small tests of your own to test this behavior. **You should not use a "brute-force" method that involves generating and/or looping over all words in the trie**.

5) Autocorrect

You may have noticed that for some words, our autocomplete implementation generates very few or no suggestions. In cases such as these, we may want to guess that the user mistyped something in the original word. We ask you to implement a more sophisticated tool: autocorrect.

In this case, we will only concern ourselves with tries that are made up of *words* (i.e., we won't concern ourselves with tuples in this case).

autocorrect(*trie, prefix, max_count*=None)

> `trie` is an instance of `Trie` whose keys are strings, `prefix` is a string, `max_count` is an integer or `None`; returns a list of up to `max_count` words. `autocorrect` should invoke `autocomplete`, but if fewer than `max_count` completions are made, suggest additional words by applying one **valid edit** to the prefix.
>
> An **edit** for a word can be any one of the following:
>
> - A single-character insertion (add any one character in the range "a" to "z" at any place in the word)
> - A single-character deletion (remove any one character from the word)
> - A single-character replacement (replace any one character in the word with a character in the range a-z)
> - A two-character transpose (switch the positions of any two adjacent characters in the word)
>
> A **valid edit** is an edit that **results in a word in the trie without considering any suffix characters**. In other words we don't try to autocomplete valid edits, we just check if `edit in trie` is True.
>
> For example, editing `"te"` to `"the"` is valid, but editing `"te"` to `"tze"` is not, as "tze" isn't a word. Likewise, editing `"phe"` to `"the"` is valid, but `"phe"` to `"pho"` is not because "pho" is not a word in the corpus, although many words beginning with "pho" are.
>
> In summary, given a prefix that produces C completions, where $C <$ `max_count`, generate up to `max_count` - C additional words by considering all valid single edits of that prefix (i.e., corpus words that can be generated by 1 edit of the original prefix) and selecting the most-frequently-occurring edited words. Return a list of suggestions produced by including **all** C of the completions and up to `max_count` - C of the most-frequently-occuring valid edits of the prefix; the list may be in any order. Be careful not to repeat suggested words!
>
> If `max_count` is `None` (or is unspecified), `autocorrect` should return all autocompletions as well as all valid edits. **You should not use a "brute-force" method that involves generating and/or looping over all words in the trie**.
>
> Example (using the example trie from above):

- `autocorrect(t, "bar", 3)` returns a list containing 'bar', 'bark', and 'bat' since "bar" and "bark" are found by autocomplete and "bat" is valid edit involving a single-character replacement, i.e., "t" is replacing the "r" in "bar".

6) Selecting words from a word trie

It's sometimes useful to select only the words from a trie that match a pattern. That's the purpose of the `filter` method.

word_filter(*trie, pattern*)

`trie` is a trie whose keys are strings, and `pattern` is a string. Return a list of `(word, freq)` tuples for those words whose characters match those of `pattern`. The characters in `pattern` are matched one at a time with the characters in each word stored in the trie. If all the characters in a particular word are matched, the `(word, freq)` pair should be included in the list to be returned. The list can be in any order.

The characters in `pattern` are interpreted as follows:

- `'*'` matches a sequence of **zero or more** of the next unmatched characters in `word`.
- `'?'` matches the next unmatched character in `word` no matter what it is. There must be a next unmatched character for `'?'` to match.
- otherwise the character in the pattern must exactly match the next unmatched character in the word.

Pattern examples:

- `"*a*t"` matches all words that contain an "a" and end in "t". This would include words like "at", "art", "saint", and "what".
- `"http://*"` matches all words that start with "http://".
- `"year*"` would match "year," "years," and "yearn," among others (as well as longer words like "yearning").
- `"year?"` would match "years" and "yearn" (but not longer words).
- `"*ing"` matches all words ending in "ing".
- `"???"` would match all 3-letter words.
- `"?ing"` matches all 4-letter words ending in "ing".
- `"?*ing"` matches all words with 4 or more letters that end in "ing".

Filter examples (using the example trie from above):

- `word_filter(t, "*")` returns a list containing the pairs `('bat', 2)`, `('bar', 1)`, and `('bark', 1)`, i.e., listing all the words in the trie.
- `word_filter(t, "???")` returns a list containing the pairs `('bat', 2)` and `('bar', 1)`, i.e., listing all the 3-letter words in the trie.
- `word_filter(t, "*r*")` returns a list containing the pairs `('bar', 1)` and `('bark', 1)`, i.e., listing all the words containing an "r" in any position.

Hint: the matching operation can implemented as a recursive search function that attempts to match the next character in the pattern with some number of characters at the beginning of the word, then recursively matches the remaining characters in the pattern with remaining unmatched characters in the word. **You should not use a "brute-force" method that involves generating and/or looping over all words in the trie.**

Note: **you cannot use any of the built-in Python pattern-matching functions**, e.g., functions from the `re` module — you are expected to write your own pattern-matching code. Copying code from StackOverflow is also not appropriate.

7) Testing your lab

As in the previous labs, we provide you with a `test.py` script to help you verify the correctness of your code. We've also included a server you can use to visualize the outputs of your `autocomplete`, `autocorrect`, and `word_filter` functions on different corpora. In addition to the test cases for this week's lab, we'll have you test out your code by running it on an example of a real public-domain book (courtesy of Project Gutenberg).

The folder `'resources/corpora'` contains text files of public-domain books that can be used as corpora for generating tries. Feel free to add text files from Project Gutenberg or elsewhere to this directory for testing. The questions below will require you to answer them using tries generated from Jane Austen's *Pride and Prejudice*.

You can load the text of a corpus file using something like the following code:

```python
with open("filename.txt", encoding="utf-8") as f:
    text = f.read()
```

After running this code, the variable `text` will be bound to a string containing the text contained in the `filename.txt` file.

We'll read the contents of these files into Python, use our `make_word_trie` and `make_phrase_trie` functions to create the relevant trie structures, and we will use our autocompletion/autocorrection based on this corpus. You can alternatively use the server interface to obtain the results of your implemented methods on the files in the `'resources/corpora'` folder.

In *Pride and Prejudice*, what are the four most common sentences (regardless of prefix)? Enter your answer as a Python list of tuples:

[   ]

This question is due on Friday October 25, 2019 at 04:00:00 PM.

In *Pride and Prejudice*, what are the six most common words starting with `gre`? Enter your answer as a Python list of strings:

[   ]

This question is due on Friday October 25, 2019 at 04:00:00 PM.

What are the top nine autocorrections for `'tear'` in *Pride and Prejudice*? Enter your answer as a Python list of strings:

[   ]

This question is due on Friday October 25, 2019 at 04:00:00 PM.

In *Pride and Prejudice*, what are all of the words matching the pattern `r?c*t`? Enter your answer as a Python list of strings:

[   ]

This question is due on Friday October 25, 2019 at 04:00:00 PM.

8) Code Submission

[Select File] No file selected
This question is due on Friday October 25, 2019 at 04:00:00 PM.

9) Checkoff

Once you are finished with the code, please come to a lab session or office hour and add yourself to the queue asking for a checkoff. **You must be ready to discuss your code in detail before asking for a checkoff.** Since the clarity of your code will be evaluated as part of the checkoff, you may wish to take some time to comment your code, use good variable names, avoid repetitive code (create helper methods), etc.

Be prepared to discuss:

- How you were able to keep track of the prefix associated with each node **without explicitly storing the prefix itself**
- The tradeoff between using iteration and recursion when implementing the `__getitem__` method.
- How your implementation of `__iter__` works.
- How using your other methods helped in implementing `autocomplete`.
- How your code for creating edits works.
- How your recursive matching works (without enumerating all words) for the `word_filter` implementation.

9.1) Grade

> *You have not yet received this checkoff. When you have completed this checkoff, you will see a grade here.*

Footnotes

[1] Different people have different opinions about whether this data structure's name should be pronounced like "tree" or like "try." It originally comes from the middle syllable of the word "retrieval," which suggests one pronunciation, but some prefer to say it like "try" to avoid confusion with general "tree" structures in programming. Some even say "tree as in try," but that's kind of a mouthful...